

# *PARALLELIZATION USING THREE DIFFERENT METHODS*

Kansas State University: Department of Computer Science

*Bristol Miller, Shreya Kumar, Corey Vessar*

*bristolnmiller@ksu.edu, shreyak@ksu.edu, coreyvessar@ksu.edu  
Manhattan, KS 66503*

## Table of Contents

<b>Abstract.....</b>	<b>2</b>
<b>Introduction.....</b>	<b>2</b>
<b>Related Work.....</b>	<b>3</b>
<b>Implementation.....</b>	<b>3</b>
<b>Evaluation .....</b>	<b>4</b>
<b>Conclusions &amp; Future Work .....</b>	<b>4</b>
<b>Appendix .....</b>	<b>5</b>
<b>OpenMP.....</b>	<b>5</b>
OpenMP Shell Script .....	8
<b>MPI .....</b>	<b>8</b>
MPI Shell Script.....	Error! Bookmark not defined.
<b>PThreads.....</b>	<b>12</b>
PThreads Shell Script .....	16
<b>References .....</b>	<b>16</b>

## Abstract

Implementing programs without parallelization can make even simple solutions quite slow. When reading a larger file and searching it line by line for information, it is necessary to parallelize your program in order to have the most efficient solution. Through our trial of the three different methods of parallelization: OpenMP, MPI, and PThreads, we have found that one method outperforms the others and improves overall timing by an average of (XXX).

## Introduction

A program that we have written reads in a 1.7 GB file, called wiki\_dump.txt, and compares the first two lines, followed by the second and third lines, then the third and fourth lines, and so on, searching for the longest common substring between each pairing and outputting those substrings to a file. The major problem with this program is that it is inefficient.

Initially, we had not implemented any parallelization in order to not complicate the code. Avoiding parallelization from the start led to our program being very slow. From start to end, accomplishing the tasks of reading in the file and searching for the longest common substrings took (XXX) amount of time. When faced with a heavy task such as reading in a large file and searching through each line of that large file, this can be detrimental if the information is needed quickly. We have solved this problem by parallelizing the program in three different ways, in order to figure out which method provides the quickest way to determine the longest common substrings. These three methods are OpenMP, MPI, and PThreads.

By creating solutions using the OpenMP, MPI, and PThreads methods all separately, we were able to discern which method is the most efficient for this application. Each of the three methods presented a considerable decrease in runtime, and thus a dramatic increase in efficiency. Also, having the comparisons between the three different methods allowed us to make an educated and unbiased decision as to which solution is the best for this specific application. After experimentation with all three methods, we came to the conclusion that the (XXX) method is the best for our program because of (X, Y, and Z).

## Related Work

One could approach this problem in a variety of ways. One such example would be to use a single-threaded implementation. This method would decrease the complexity of such a solution, but the solution would be much slower than our implementations. Usually when you convert a program from single-threaded to multi-threaded, you get a dramatic speedup without too much effort. There is, of course, a balance to strike between the two, though. For some applications, the speed is less important, so it is considered a waste of time to turn the program into a multi-threaded one. On the other hand, there are plenty of scenarios where time is of the essence, and a faster program is valued higher than the effort required to make it so.

Another way to vary the approach of this problem would be to copy the file locally and use that instead of reading it in to memory from a separate location. Approaching with this manner isn't very spatially effective, but it may produce a speedup as far as timing is concerned. On this note, it also might be more effective to parallelize reading the file in. In our implementation, we simply parallelized the search, rather than our entire program. If we were to read multiple lines into memory at a time, we're certain that this would be quicker, although it would be more difficult to code.

There are multiple different methods one can use to parallelize their solution. The three that we decided to test were OpenMP, PThreads, and MPI, however, there are more available to programmers. Some of those other methods include OpenACC, OpenCL, and NVIDIA CUDA. For our purposes, we aren't considering the other methods, although they have their own advantages and drawbacks. CUDA tends to have very high performance, but often requires frequent updates and is expensive. OpenCL may not have as high of performance, but it isn't very expensive and doesn't require frequent updates.

## Implementation

MPI is a method that implements message passing, in fact, it's name is an acronym for Message Passing Interface. In message passing, there are two main functions used: send and receive. Send allows threads communicate with each other. This can be either synchronous or asynchronous. If its synchronous, then it waits for the

receive end to send an acknowledgement. On the other hand, if its asynchronous, the message passed from one thread to another can get pushed into a queue, and the thread that sent it can continue about its business. When a thread reaches the receive message, it usually waits, unless there is a message in the queue. When a message is received, we send an acknowledgement.

OpenMP is a different method in which part of your program is single-threaded, while other parts are multi-threaded. The M and P in OpenMP are short for Multi-Processing. This method overall has three different components: a runtime library, environment variables that define the runtime parallel parameters, and directives for the programs. OpenMP is known for its ease of use and portability, while allowing users to standardize and “lean-out” their code.

PThreads is short for POSIX Threads. It is a method of parallel programming in which the threads are incredibly light-weight and efficient with data exchange. Within this method, there are four main types of functions included in its library. These types are mutexes, condition variables, synchronization between threads using locks and barriers on reads and writes, and thread management techniques including creating and joining threads.

## Evaluation

Our methods for testing involved submitting a variety of jobs to Beocat, Kansas State University’s supercomputer. We chose to run the three different methods, along with our single-threaded implementation as a control, with a combination of (X,Y, and Z) cores with a job size varying from 1000 lines to 1 million lines for comparison. Through these evaluation steps we were able to find that for fewer lines (XXX) performed best, while for more lines (YYY) performed best, and overall the parallelized methods out-performed the single-threaded method.

## Conclusions & Future Work

We have found that out of the following three methods of parallelization: OpenMP, MPI, and PThreads, (XXX) is the best for our application because of (X, Y, and Z). Thus, when it is necessary to read in a file and search that file for the longest common substrings, we have found one of the most efficient ways to solve the problem.

In the future, one could compare other methods to our findings in this paper. Another way to improve upon our work would be to parallelize the reads, although this would provide only slightly better timing results.

## Appendix

In this section, we have included copies of our code, along with links to the files on github.

### OpenMP

<https://github.com/shreyakumar1010/cis520-Project4/blob/master/OpenMP/OpenMP.c>

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/time.h>
#include <stdbool.h>
#include <omp.h>
#define WIKI_ARRAY_SIZE 50000
#define WIKI_LINE_SIZE 2001
#define num_threads 32
int LCS(char * s1, char * s2, char ** longest_common_substring);
//load the lines into an array
char **wiki_array;
char **longestSub;
omp_lock_t my_lock;
void readToMemory();
void printResults();
void printToFile();
int main()
{
    omp_init_lock(&my_lock);
    struct timeval time1;
    struct timeval time2;
    struct timeval time3;
    struct timeval time4;
    double e1, e2, e3;
    int numSlots, Version = 3; //base = 1, pthread = 2, openmp = 3, mpi = 4
    gettimeofday(&time1, NULL);
    readToMemory();
    gettimeofday(&time2, NULL);
    //time to read to memory
    e1 = (time2.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
    e1 += (time2.tv_usec - time1.tv_usec) / 1000.0; // us to ms
    printf("Time to read full file to Memory: %f\n", e1);
    gettimeofday(&time3, NULL);
    int i,j, startPos, endPos, myID;
    char** longestSubPointer;
    omp_set_num_threads(num_threads);

    #pragma omp parallel private(myID, startPos, endPos, j, longestSubPointer)
    {
        myID = omp_get_thread_num();
        startPos = (myID) * (WIKI_ARRAY_SIZE / num_threads);
        endPos = startPos + (WIKI_ARRAY_SIZE / num_threads);
        longestSubPointer = longestSub + startPos;
    }
    if(myID == num_threads-1)
    {
```

```

        endPos = WIKI_ARRAY_SIZE - 1 ;
    }
        for (j = startPos; j < endPos; j++)
        {
            LCS((void*)wiki_array[j], (void*)wiki_array[j+1], longestSubPointer);
            longestSubPointer++;
        }
    }
    printToFile();
    gettimeofday(&time4, NULL);
    //time to find all longest substrings
    e2 = (time4.tv_sec - time3.tv_sec) * 1000.0; //sec to ms
    e2 += (time4.tv_usec - time3.tv_usec) / 1000.0; // us to ms
    printf("Time find all Substrings: %f\n", e2);
    //total elapsed time between reading and finding all longest substrings
    e3 = (time4.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
    e3 += (time4.tv_usec - time1.tv_usec) / 1000.0; // us to ms
    printf("DATA, %d, %s, %f, %d\n", myVersion, getenv("NSLOTS"), e3, num_threads);
}
void readToMemory()
{
    int nlines, maxlines = 10;
    int k, n, err, *count;
    int i;
    double nchars = 0;
    FILE *fd;
    //Adding malloc for space
    wiki_array = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));
    for (i = 0; i < WIKI_ARRAY_SIZE; i++)
    {
        wiki_array[i] = malloc(2001);
    }
    //saved results
    longestSub = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));
    for (i = 0; i < WIKI_ARRAY_SIZE -1; i++)
    {
        longestSub[i] = malloc(2001);
    }
    fd = fopen("/homes/dan/625/wiki_dump.txt", "r");
    nlines = -1;
    do
    {
        err = fscanf(fd, "%[^\n]\n", wiki_array[++nlines]);
        if(wiki_array[nlines] != NULL)
            nchars += (double) strlen(wiki_array[nlines]);
    }
    while (err != EOF && nlines < WIKI_ARRAY_SIZE);

    fclose(fd);
    printf("Read in %d lines averaging %.01f chars/line\n", nlines, nchars / nlines);
}

```

```

void printToFile()
{

```

```

FILE *f = fopen("LargestCommonSubstringsONETHREAD.txt", "w");
if (f == NULL)
{
    printf("Error opening LargestCommonSubstrings.txt!\n");
    exit(1);
}
int i;
for(i = 0; i < WIKI_ARRAY_SIZE - 2; i++)
{
    fprintf(f, "%d-%d: %s", i, i + 1, longestSub[i]);
    fprintf(f, "\n");
}
fclose(f);
}
void printResults()
{
    int i;
    for(i = 0; i <= WIKI_ARRAY_SIZE - 2; i++)
    {
        printf("%d-%d: %s", i, i + 1, longestSub[i]);
        printf("\n");
    }
}
int LCS(char *s1, char *s2, char **longest_common_substring)
{
    int s1_length = strlen(s1);
    int s2_length = strlen(s2);

    int **_matrix;
    int _matrix_row_size = 0;
    int _matrix_column_size = 0;
    if (s1_length+1 > _matrix_row_size || s2_length+1 > _matrix_column_size)
    {
        int i;
        /* malloc matrix */
        _matrix = (int **)malloc((s1_length+1) * sizeof(int*));
        for (i = 0; i < s1_length+1; i++)
            _matrix[i] = (int *)malloc((s2_length+1) * sizeof(int));
        _matrix_row_size = s1_length+1;
        _matrix_column_size = s2_length+1;
    }
    int i;
    for (i = 0; i <= s1_length; i++)
        _matrix[i][s2_length] = 0;
    int j;
    for (j = 0; j <= s2_length; j++)
        _matrix[s1_length][j] = 0;
    int max_len = 0, max_index_i = -1;

    for (i = s1_length-1; i >= 0; i--)
    {
        for (j = s2_length-1; j >= 0; j--)
        {
            if (s1[i] != s2[j])
            {
                _matrix[i][j] = 0;
                continue;
            }
            _matrix[i][j] = _matrix[i+1][j+1] + 1;

            if (_matrix[i][j] > max_len)
            {

```



```

        max_len = _matrix[i][j];
        max_index_i = i;
    }
}
if (longest_common_substring != NULL)
{
    omp_set_lock(&my_lock);
    *longest_common_substring = malloc(sizeof(char) * (max_len+1));
    strncpy(*longest_common_substring, s1+max_index_i, max_len);
    (*longest_common_substring)[max_len] = '\0';
    omp_unset_lock(&my_lock);
}
/* free matrix */
for (i = 0; i < _matrix_row_size; i++)
    free(_matrix[i]);
free(_matrix);
return max_len;
}

```

## OpenMP Shell Script

<https://github.com/shreyakumar1010/cis520-Project4/blob/master/OpenMP/openmp.sh>

```

##### These lines are for Slurm
#SBATCH -N 4          #Number of nodes to use
#SBATCH -p pReserved  #Workshop queue
#SBATCH -t 5:00       #Maximum time required
#SBATCH -o output.%j  #Output file name

### Job commands start here
### Display some diagnostic information
echo '=====JOB DIAGNOTICS=====
date
echo -n 'This machine is ';hostname
echo -n 'My jobid is '; echo $SLURM_JOBID
echo 'My path is:'
echo $PATH
echo 'My job info:'
squeue -j $SLURM_JOBID
echo 'Machine info'
sinfo -s

echo '=====JOB STARTING=====

### CHANGE THE LINES BELOW TO SELECT DIFFERENT MPI CODES AND/OR COMPILERS
#Compile an exercise code
gcc -fopenmp OpenMP.c -o OpenMP
#Run the code
srun -N4 ./OpenMP

### Issue the sleep so we have time to see the job actually running
sleep 120
echo ''
echo '=====ALL DONE=====

```

## MPI

<https://github.com/shreyakumar1010/cis520-Project4/blob/master/MPI/MPI.c>

```
#include <stdio.h>
```

```

#include <stdlib.h>
#include <string.h>
#include <sys/time.h>
#include <stdbool.h>
#include <mpi.h>

#define WIKI_ARRAY_SIZE 500
#define WIKI_LINE_SIZE 2001
#define MASTER 0

//int lengthOfSubstring [WIKI_ARRAY_SIZE];
int LCS (char * s1, char * s2, char ** longest_common_substring);
int num_threads;
char * numCores = "1";

char ** wiki_array;
char ** longestSub;
char ** localLongestSub;

void readToMemory();
void printResults();
void printToFile();
void * findem(void * rank);

int main(int argc, char* argv[])
{
    struct timeval time1;
    struct timeval time2;
    struct timeval time3;
    struct timeval time4;
    double e1, e2, e3;
    int numSlots, Version = 4; //base = 1, pthread = 2, openmp = 3, mpi = 4
    //longestSub = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *)); //saved results
    int i;
    //for (i = 0; i < WIKI_ARRAY_SIZE -1; i++)//initializing saved results
    //    longestSub[i] = malloc(2001);

    //=====INIT MPI=====
    int rc, NumTasks, rank;
    MPI_Status status;
    rc = MPI_Init(&argc, &argv);
    if(rc != MPI_SUCCESS)
    {
        printf("MPI isn't working");
        MPI_Abort(MPI_COMM_WORLD, rc);
    }
    MPI_Comm_size(MPI_COMM_WORLD, &NumTasks);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    num_threads = NumTasks;

    //=====READ TO MEMORY=====
    gettimeofday(&time1, NULL);

    <<<<<<< HEAD
    readToMemory(); //reading
    =====
    gettimeofday(&time2, NULL); //time to read to memory
    e1 = (time2.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
    e1 += (time2.tv_usec - time1.tv_usec) / 1000.0; // us to ms
    printf("Time to read full file to Memory: %f\n", e1);

    //====FINDING LONGEST CMN SUBSTR. & PARALLELIZING====
    >>>>>> 75b7374bf3009e45a76c1b5ab6a6927162442560

```

```

gettimeofday(&time3, NULL);
MPI_Bcast(wiki_array, WIKI_ARRAY_SIZE * WIKI_LINE_SIZE, MPI_CHAR, o, MPI_COMM_WORLD);
findem(&rank);

MPI_Reduce(localLongestSub, longestSub, WIKI_ARRAY_SIZE * WIKI_LINE_SIZE, MPI_CHAR,
MPI_SUM, o, MPI_COMM_WORLD);
if(rank == MASTER)
{
//=====FINAL TIMING=====
gettimeofday(&time4, NULL);

printToFile();
//time to find all longest substrings
e2 = (time4.tv_sec - time3.tv_sec) * 1000.0; //sec to ms
e2 += (time4.tv_usec - time3.tv_usec) / 1000.0; // us to ms
printf("Time find all Substrings: %f\n", e2);
//total elapsed time between reading and finding all longest substrings
e3 = (time4.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
e3 += (time4.tv_usec - time1.tv_usec) / 1000.0; // us to ms
printf("DATA, %d, %s, %f, %d, %d, %s\n", Version, getenv("NSLOTS"), e3, num_threads,
WIKI_ARRAY_SIZE, numCores);
}
<<<<<<< HEAD
else {
MPI_Finalize();
return -1;
}

MPI_Finalize();
return 0;
=====
MPI_Finalize();
>>>>>>> 75b7374bf3009e45a76c1b5ab6a6927162442560
}

void * findem(void * rank)
{
int myID = *((int *) rank);
int startPos = ((long) myID) * (WIKI_ARRAY_SIZE / num_threads);
int endPos = startPos + (WIKI_ARRAY_SIZE / num_threads);
if(myID == num_threads - 1)
endPos = WIKI_ARRAY_SIZE;
int i, j;

for(j = startPos; j < endPos; j++)
{
LCS((void*)wiki_array[i], (void*)wiki_array[i+1], localLongestSub);
localLongestSub++;
}
}

void readToMemory()
{
int nlines, maxlines = 10;
int k, n, err, *count, nthreads = 24;
int i;
double nchars = 0;
FILE *fd;
wiki_array = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));
for (i = 0; i < WIKI_ARRAY_SIZE; i++)//initializing the array
{
wiki_array[i] = malloc(2001);

```

```

    }
    longestSub = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));
    localLongestSub = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));
    //=====READING FILE=====
    fd = fopen("/homes/dan/625/wiki_dump.txt", "r");
    nlines = -1;
    do
    {
        err = fscanf(fd, "%[^\n]\n", wiki_array[++nlines]);
        if(wiki_array[nlines] != NULL)
            nchars += (double) strlen(wiki_array[nlines]);
    }
    while (err != EOF && nlines < WIKI_ARRAY_SIZE);
    fclose(fd);
    printf("Read in %d lines averaging %.01f chars/line\n", nlines, nchars / nlines);
}

void printToFile()//self-explanatory. Prints the substrings found to a file for processing
{
    FILE *f = fopen("LargestCommonSubstrings.txt", "w");
    if (f == NULL)//if there's a problem with the file error handling
    {
        printf("Error opening LargestCommonSubstrings.txt!\n");
        exit(1);
    }
    //longestSub = longestSub - (WIKI_ARRAY_SIZE - 1);
    int i;
    for(i = 0; i < WIKI_ARRAY_SIZE - 2; i++)
    {
        fprintf(f, "%d-%d: %s", i, i + 1, longestSub[i]);
        fprintf(f, "\n");
    }
    fclose(f);
}

void printResults()//printing to console for debugging
{
    int i;
    //longestSub = longestSub - (WIKI_ARRAY_SIZE - 1);
    for(i = 0; i <= WIKI_ARRAY_SIZE - 2; i++)
    {
        printf("%d-%d: %s", i, i + 1, longestSub[i]);
        printf("\n");
    }
}

int LCS(char *s1, char *s2, char **longest_common_substring)
{
    int s1_length = strlen(s1);
    int s2_length = strlen(s2);
    int **_matrix;
    int _matrix_row_size = 0;
    int _matrix_column_size = 0;
    if (s1_length+1 > _matrix_row_size || s2_length+1 > _matrix_column_size)
    {
        int i;
        /* malloc matrix */
        _matrix = (int **) malloc((s1_length+1) * sizeof(int*));
        for (i = 0; i < s1_length+1; i++)
            _matrix[i] = (int *) malloc((s2_length+1) * sizeof(int));
        _matrix_row_size = s1_length+1;
        _matrix_column_size = s2_length+1;
    }
}

```

```

int i;
for (i = 0; i <= s1_length; i++)
    _matrix[i][s2_length] = 0;
int j;
for (j = 0; j <= s2_length; j++)
    _matrix[s1_length][j] = 0;
int max_len = 0, max_index_i = -1;
for (i = s1_length-1; i >= 0; i--)
{
    for (j = s2_length-1; j >= 0; j--)
    {
        if (s1[i] != s2[j])
        {
            _matrix[i][j] = 0;
            continue;
        }
        _matrix[i][j] = _matrix[i+1][j+1] + 1;
        if (_matrix[i][j] > max_len)
        {
            max_len = _matrix[i][j];
            max_index_i = i;
        }
    }
}
if (longest_common_substring != NULL)
{
    *longest_common_substring = malloc(sizeof(char) * (max_len+1));
    strncpy(*longest_common_substring, s1+max_index_i, max_len);
    (*longest_common_substring)[max_len] = '\0';
}
/* free matrix */
for (i = 0; i < _matrix_row_size; i++)
    free(_matrix[i]);
free(_matrix);
return max_len;
}

```

## PThreads

<https://github.com/shreyakumar1010/cis520-Project4/blob/master/Pthreads/Pthreads.c>

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/time.h>
#include <stdbool.h>
#include <pthread.h>

#define WIKI_ARRAY_SIZE 50000
#define WIKI_LINE_SIZE 2001

int lengthOfSubstring [WIKI_ARRAY_SIZE];
int LCS (char * s1, char * s2, char ** longest_common_substring);
void loopingFunc(void *myID);
char * numCores = "0";

int num_threads = 32;

//load the lines into an array
char **wiki_array;
char **longestSub;

void readToMemory();

```

```

void printResults();
void printToFile();

int main()
{
    struct timeval time1;
    struct timeval time2;
    struct timeval time3;
    struct timeval time4;
    double e1, e2, e3;
    int numSlots, Version = 2; //base = 1, pthread = 2, openmp = 3, mpi = 4

    int myID, k, rc, i;
    pthread_t threads[num_threads];
    pthread_attr_t attr;
    void *status;

    /* Initialize and set thread detached attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

    gettimeofday(&time1, NULL);
    readToMemory();
    gettimeofday(&time2, NULL);

    //time to read to memory
    e1 = (time2.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
    e1 += (time2.tv_usec - time1.tv_usec) / 1000.0; // us to ms
    printf("Time to read full file to Memory: %f\n", e1);

    gettimeofday(&time3, NULL);

    for (i = 0; i < num_threads; i++)
    {
        rc = pthread_create(&threads[i], &attr, loopingFunc, (void *)i);
        if(rc)
        {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    pthread_attr_destroy(&attr);

    for(i=0; i<num_threads; i++)
    {
        rc = pthread_join(threads[i], &status);
        if (rc)
        {
            printf("ERROR; return code from pthread_join() is %d\n", rc);
            exit(-1);
        }
    }

    //printResults();
    printToFile();

    gettimeofday(&time4, NULL);

    //time to find all longest substrings
    e2 = (time4.tv_sec - time3.tv_sec) * 1000.0; //sec to ms
    e2 += (time4.tv_usec - time3.tv_usec) / 1000.0; // us to ms

```

```

    printf("Time find all Substrings: %f\n", e2);

    //total elapsed time between reading and finding all longest substrings
    e3 = (time4.tv_sec - time1.tv_sec) * 1000.0; //sec to ms
    e3 += (time4.tv_usec - time1.tv_usec) / 1000.0; // us to ms
    printf("DATA, %d, %s, %f, %d, %d, %s\n", Version, getenv("NSLOTS"), e3, num_threads,
WIKI_ARRAY_SIZE, numCores);

}

void loopingFunc(void *myID)
{
    //start position of the array
    int startPos = ((int) myID) * (WIKI_ARRAY_SIZE / num_threads);
    char ** tempPos = longestSub + startPos;
    //end position of the array
    int endPos = startPos + (WIKI_ARRAY_SIZE / num_threads);
    //longestSub = longestSub + startPos;

    if((int)myID == num_threads -1)
    {
        endPos = WIKI_ARRAY_SIZE - 1;
    }

    int i, j;

    //for(i = 0; i < WIKI_ARRAY_SIZE - 1; i++)
    //{
        for(j = startPos; j < endPos; j++)
        {
            printf("%d-%d: %s", j, j + 1, "lines submitted to LCS");
            printf("\n");
            LCS((void*)wiki_array[j], (void*)wiki_array[j+1], tempPos);
            tempPos++;
        }
    //}
    pthread_exit(NULL);
}

void readToMemory()
{
    int nlines, maxlines = 10;
    int k, n, err;
    int i;
    double nchars = 0;
    FILE *fd;

    //Adding malloc for space
    wiki_array = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));

    for (i = 0; i < WIKI_ARRAY_SIZE; i++)
    {
        wiki_array[i] = malloc(2001 * sizeof(char));
    }
    //saved results
    longestSub = (char **) malloc( WIKI_ARRAY_SIZE * sizeof(char *));

    /*for (i = 0; i < WIKI_ARRAY_SIZE -1; i++)
    {
        longestSub[i] = malloc(2001 * sizeof(char));
    }*/
}

```

```

    fd = fopen("/homes/dan/625/wiki_dump.txt", "r");
    //fd = fopen("/homes/coreyvessar/cis520/cis520-Project4/wiki_dump.txt", "r");
    nlines = -1;
    do
    {
        err = fscanf(fd, "%[^\\n]\\n", wiki_array[++nlines]);
        if(wiki_array[nlines] != NULL)
            nchars += (double) strlen(wiki_array[nlines]);
    }
    while (err != EOF && nlines < WIKI_ARRAY_SIZE);

    fclose(fd);
    printf("Read in %d lines averaging %.01f chars/line\\n", nlines, nchars / nlines);
}

void printToFile()
{
    FILE *f = fopen("threadsCommonSubstrings.txt", "w");
    if (f == NULL)
    {
        printf("Error opening LargestCommonSubstrings.txt!\\n");
        exit(1);
    }

    //longestSub = longestSub - (WIKI_ARRAY_SIZE - 1);
    int i;
    for(i = 0; i < WIKI_ARRAY_SIZE - 2; i++)
    {
        fprintf(f, "%d-%d: %s", i, i + 1, longestSub[i]);
        fprintf(f, "\\n");
    }

    fclose(f);
}

void printResults()
{
    int i;
    //longestSub = longestSub - (WIKI_ARRAY_SIZE - 1);
    for(i = 0; i <= WIKI_ARRAY_SIZE - 2; i++)
    {
        printf("%d-%d: %s", i, i + 1, longestSub[i]);
        printf("\\n");
    }
}

int LCS(char *s1, char *s2, char **longest_common_substring)
{
    int s1_length = strlen(s1);
    int s2_length = strlen(s2);

    int **_matrix;
    int _matrix_row_size = 0;
    int _matrix_collumn_size = 0;

    if (s1_length+1 > _matrix_row_size || s2_length+1 > _matrix_collumn_size)
    {
        int i;
        /* malloc matrix */
        _matrix = (int **)malloc((s1_length+1) * sizeof(int*));
        for (i = 0; i < s1_length+1; i++)

```



```

        _matrix[i] = (int *)malloc((s2_length+1) * sizeof(int));
        _matrix_row_size = s1_length+1;
        _matrix_collumn_size = s2_length+1;
    }

    int i;
    for (i = 0; i <= s1_length; i++)
        _matrix[i][s2_length] = 0;

    int j;
    for (j = 0; j <= s2_length; j++)
        _matrix[s1_length][j] = 0;

    int max_len = 0, max_index_i = -1;
    for (i = s1_length-1; i >= 0; i--)
    {
        for (j = s2_length-1; j >= 0; j--)
        {
            if (s1[i] != s2[j])
            {
                _matrix[i][j] = 0;
                continue;
            }

            _matrix[i][j] = _matrix[i+1][j+1] + 1;

            if (_matrix[i][j] > max_len)
            {
                max_len = _matrix[i][j];
                max_index_i = i;
            }
        }
    }

    if (longest_common_substring != NULL)
    {
        //omp_set_lock(&theLock);
        *longest_common_substring = malloc(sizeof(char) * (max_len+1));
        strncpy(*longest_common_substring, s1+max_index_i, max_len);
        (*longest_common_substring)[max_len] = '\0';
        //omp_unset_lock(&theLock);
        //printf("%s\n", *longest_common_substring);
    }
    /* free matrix */
    for (i = 0; i < _matrix_row_size; i++)
        free(_matrix[i]);
    free(_matrix);
    return max_len;
}

```

## PThreads Shell Script

<https://github.com/shreyakumar1010/cis520-Project4/blob/master/Pthreads/pthread.sh>

##### These lines are for Slurm

```

#SBATCH -N 2          #Number of nodes to use

#SBATCH -p pReserved  #Workshop queue

#SBATCH -t 5:00       #Maximum time required

#SBATCH -o output.%j   #Output file name

### Job commands start here

### Display some diagnostic information

echo '=====JOB DIAGNOTICS===== '

date

echo -n 'This machine is ';hostname

echo -n 'My jobid is '; echo $SLURM_JOBID

echo 'My path is:'

echo $PATH

echo 'My job info:'

squeue -j $SLURM_JOBID

echo 'Machine info'

sinfo -s


echo '=====JOB STARTING===== '


### CHANGE THE LINES BELOW TO SELECT DIFFERENT MPI CODES AND/OR COMPILERS

#Compile an exercise code

gcc -lpthread pthread.c

#Run the code

srun -N2 ./a.out

### Issue the sleep so we have time to see the job actually running

sleep 120

echo ''

echo '=====ALL DONE===== '

```

## References

- [1] Hughes, Christian J. “ChristianJHughes/Pintos-project4.” GitHub, 9 May 2017, [github.com/ChristianJHughes/pintos-project4](https://github.com/ChristianJHughes/pintos-project4).
- [2] Barney, Blaise. “Open MP.” OpenMP, U.S. Department of Energy, 17 July 2017, [computing.llnl.gov/tutorials/openMP/](https://computing.llnl.gov/tutorials/openMP/).
- [3] Barney, Blaise. “POSIX Threads Programming.” POSIX Threads Programming, U.S. Department of Energy, 7 Mar. 2017, [computing.llnl.gov/tutorials/pthreads/](https://computing.llnl.gov/tutorials/pthreads/).
- [4] Barney, Blaise. “Message Passing Interface (MPI).” Message Passing Interface (MPI), U.S. Department of Energy, 19 June 2017, [computing.llnl.gov/tutorials/mpi/](https://computing.llnl.gov/tutorials/mpi/).
- [5] “Longest Common Substring Problem.” Wikipedia, Wikimedia Foundation, 9 Nov. 2017, [en.wikipedia.org/wiki/Longest\\_common\\_substring\\_problem](https://en.wikipedia.org/wiki/Longest_common_substring_problem).
- [6] Hancy. “Hancyxhx/Longest-Common-Substring.” GitHub, 20 Mar. 2014, [github.com/hancyxhx/Longest-Common-Substring](https://github.com/hancyxhx/Longest-Common-Substring).
- [7] “Parallel Programming Techniques.” Parallel Programming Techniques, 27 Mar. 2017, [jcsites.juniata.edu/faculty/rhodes/smui/parex.htm](http://jcsites.juniata.edu/faculty/rhodes/smui/parex.htm).
- [8] “What Is OpenMP.” Into To OpenMP, Dartmouth College, 23 Mar. 2009, [www.dartmouth.edu/~rc/classes/intro\\_openmp/](http://www.dartmouth.edu/~rc/classes/intro_openmp/).
- [9] “POSIX Threads.” Wikipedia, Wikimedia Foundation, 24 Apr. 2018, [en.wikipedia.org/wiki/POSIX\\_Threads](https://en.wikipedia.org/wiki/POSIX_Threads).
- [10] Schroeder, Simon. “Which One Do You Prefer: UDA or OpenCL.” Research Gate, Research Gate, 2014, [www.researchgate.net/post/Which\\_one\\_do\\_you\\_prefer\\_CUDA\\_or\\_OpenCL](https://www.researchgate.net/post/Which_one_do_you_prefer_CUDA_or_OpenCL).

[11]